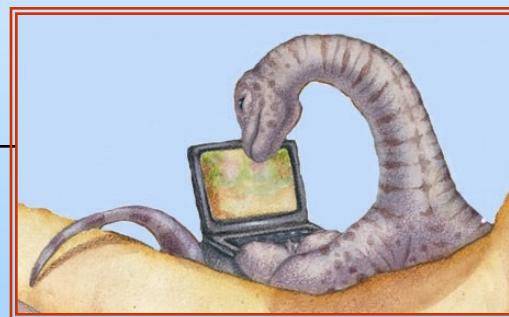




The picture can't be displayed.

# Chapter 3: Processes



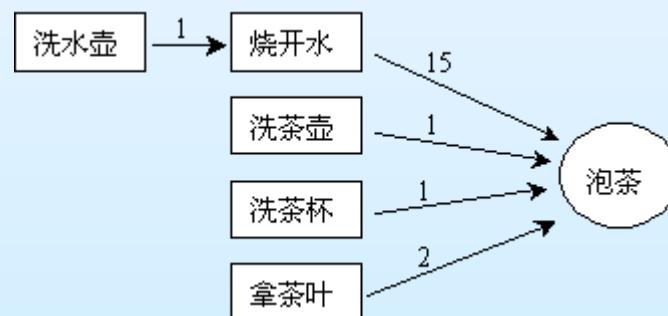


# A Lesson from High School

比如，想泡壶茶喝。当时的情况是：开水没有。开水壶要洗，茶壶茶杯要洗；火已升了，茶叶也有了。怎么办？

1. 办法甲：洗好开水壶，灌上凉水，放在火上；在等待水开的时候，洗茶壶、洗茶杯、拿茶叶；等水开了，泡茶喝。
2. 办法乙：先做好一些准备工作，洗开水壶，洗壶杯，拿茶叶；一切就绪，灌水烧水；坐待水开了，泡茶喝。
3. 办法丙：洗净开水壶，灌上凉水，放在火上；坐待水开，开了之后急急忙忙找茶叶，洗壶杯，泡茶喝。

哪一种办法省时间？谁都能一眼看出，第一种办法好，因为后二种办法都“窝了工”





# Chapter 3: Processes

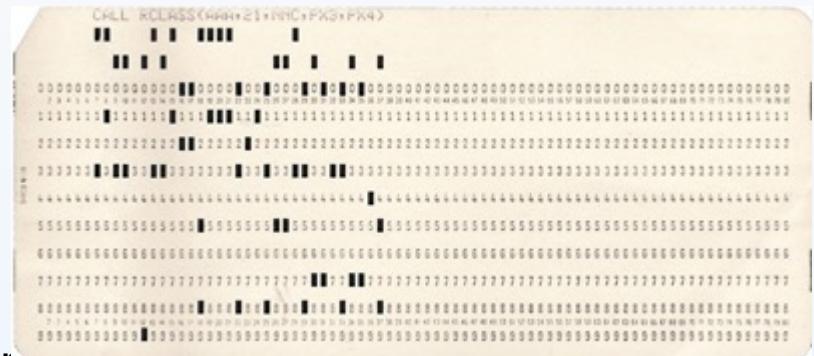
- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems





# Process Concept

- An operating system executes a variety of programs:
  - Batch system – jobs



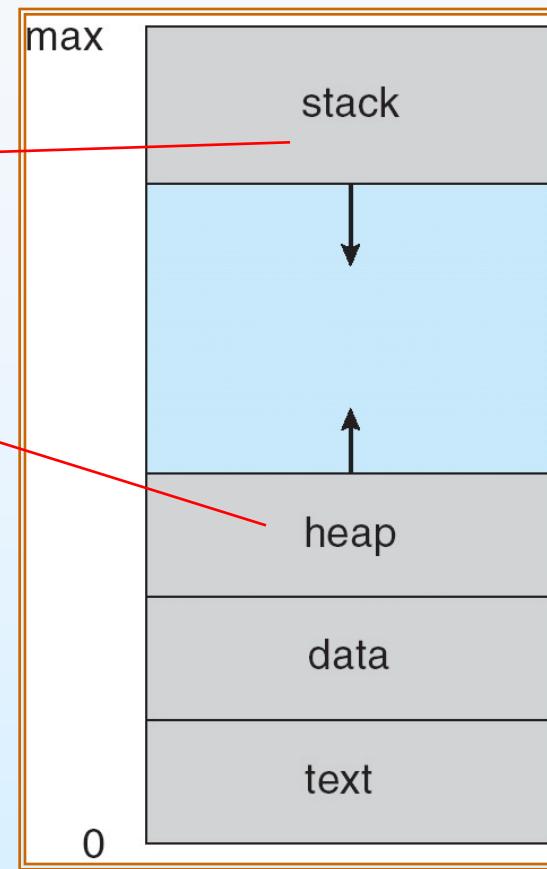
- Time-shared systems – user programs or tasks
- In textbook, *job* = *process*
- Process – a program in execution using sequential fashion
- A process includes:
  - text section (code)
  - program counter
  - stack (function parameters, local vars, return addresses)
  - data section (global vars)
  - heap (dynamically allocated memory)





# Process in Memory

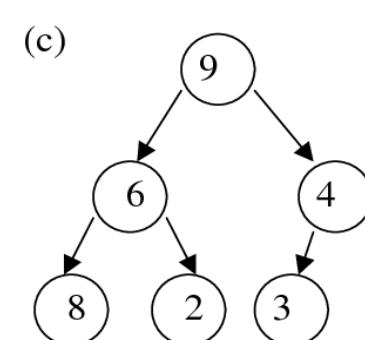
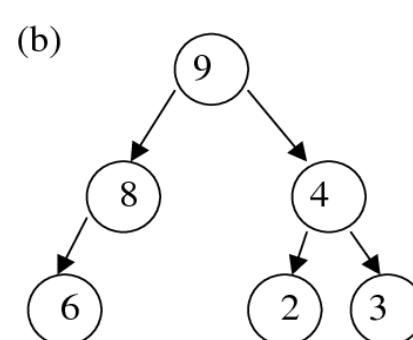
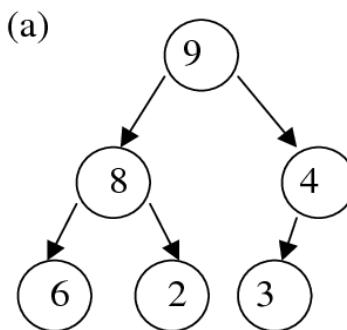
```
class A{
public:
    const int a = 0;
    ...
    int getValue(){
        int[] buffer = new int[1024];
        ...
        delete buffer;
        return 0;
    }
    ...
}
```





# Data Structure: Heap

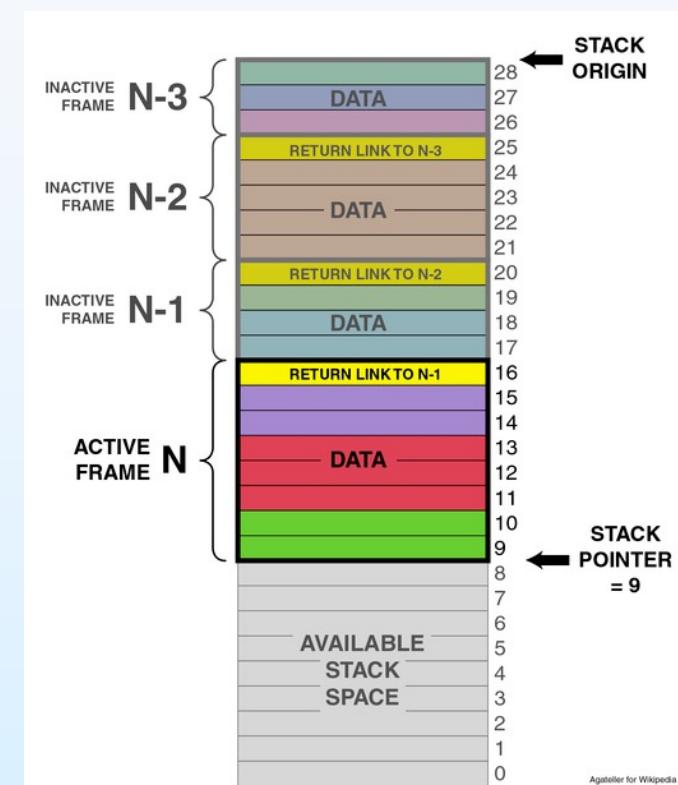
- Heap is a "complete" binary tree
  - Except the leaf level, all levels are full. The leaf level is filled from left to right
  - The node's value should be larger/smaller than its children
- (a) is a heap
- (b) is not for it is not complete,
- (c) is not because  $6 < 8$





# Data Structure: Stack

- A typical stack is an area of computer memory with a fixed origin and a variable size.
- Initially the size of the stack is zero.
- A stack pointer, usually in the form of a hardware register, points to the most recently referenced location on the stack
- a push operation, in which a data item is placed at the location pointed to by the stack pointer, and the address in the stack pointer is adjusted by the size of the data item;
- a pop or pull operation: a data item at the current location pointed to by the stack pointer is removed, and the stack pointer is adjusted by the size of the data item.





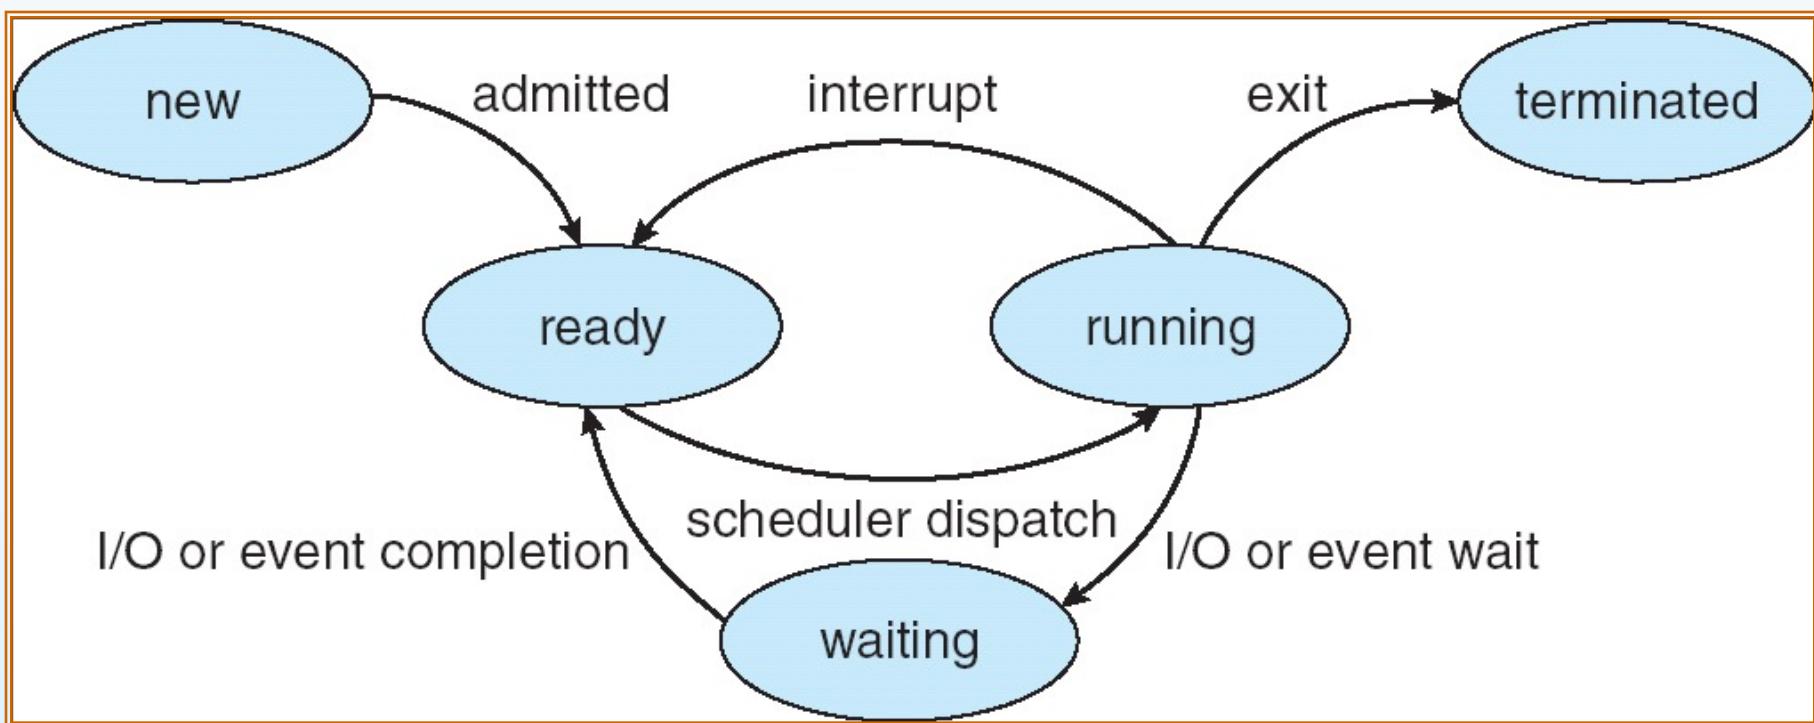
# Process State

- As a process executes, it changes *state*
  - **new**: The process is being created
  - **running**: Instructions are being executed
  - **waiting**: The process is waiting/blocked for some event to occur
  - **ready**: The process is waiting to be assigned to a processor
  - **terminated**: The process has finished execution





# Diagram of Process State





# Process Control Block (PCB)

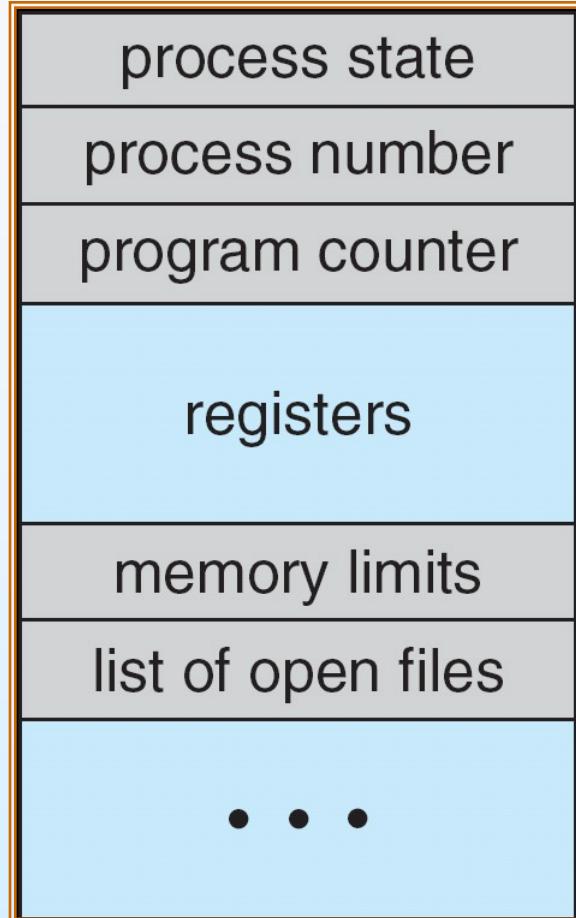
Information associated with each process

- Process state (new, ready, ...)
- Program counter (address of next instruction)
- Contents of CPU registers
- CPU scheduling information (priority)
- Memory-management information
- Accounting information (cpu/time used, time limits, process number)
- I/O status information (allocated devices and opened files)





# Process Control Block (PCB)



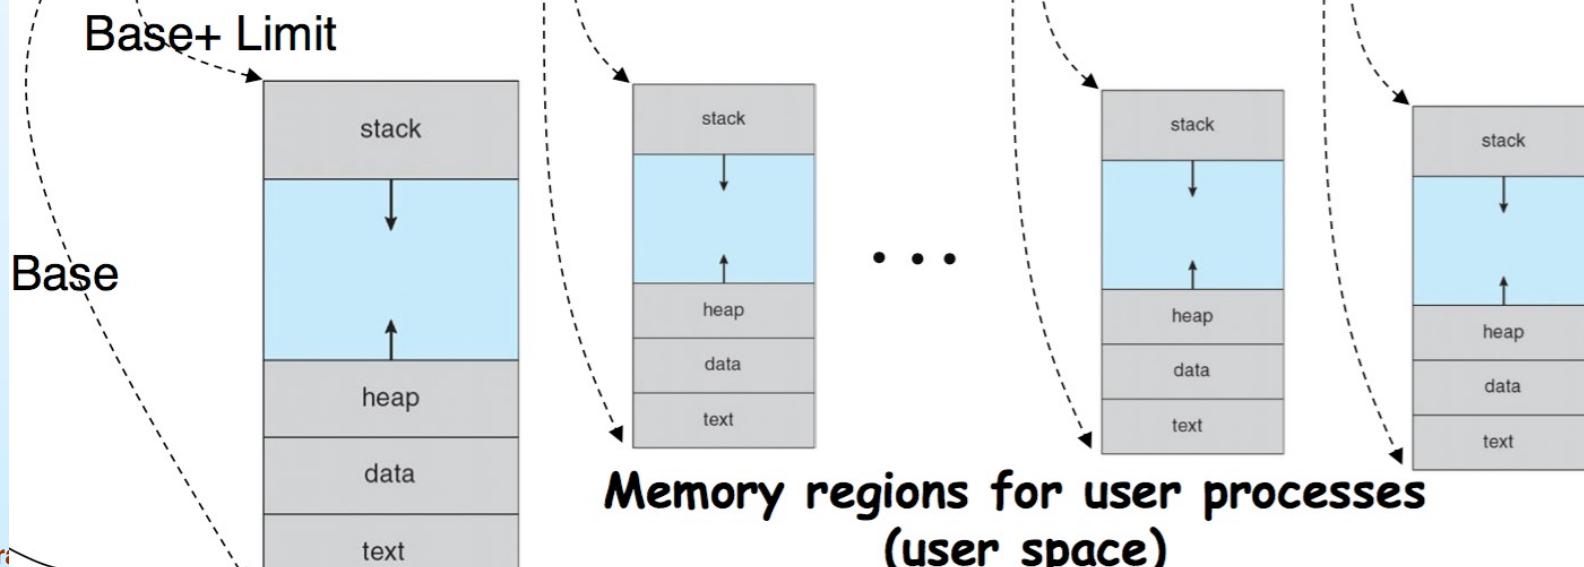
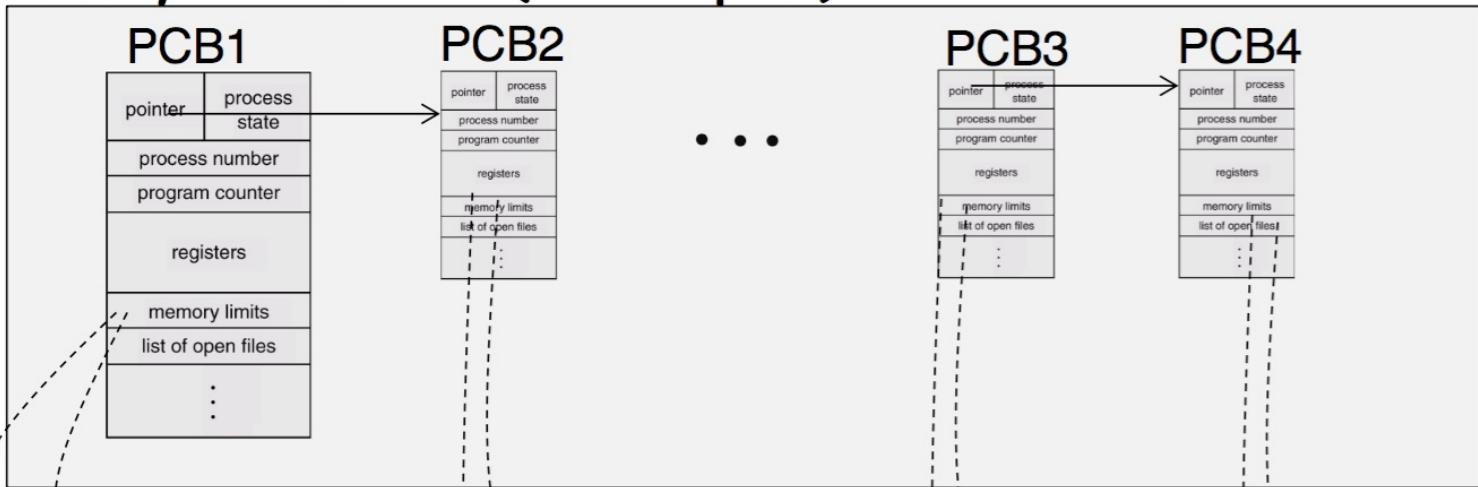
F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	512	2667	32203	0	80	0	-	26525	wait	pts/12	00:00:00	sh
0	D	512	2668	2667	29	80	0	-	50494499	sync_p	pts/12	01:56:34	convert_imagese
0	S	504	4923	27549	0	80	0	-	39358	n_tty_	pts/21	00:00:00	mysql
0	S	529	6508	6355	0	80	0	-	3894	poll_s	pts/6	00:00:16	top
0	R	501	6800	6773	0	80	0	-	27033	-	pts/32	00:00:00	ps
0	S	523	7485	13843	0	80	0	-	966037	futex_	pts/15	03:48:24	MATLAB
0	S	523	10629	1194	0	80	0	-	15047	n_tty_	pts/8	00:00:00	ssh
0	S	503	14152	14135	0	80	0	-	26524	wait	pts/5	00:00:00	nsight
0	S	503	14156	14152	0	80	0	-	26555	wait	pts/5	00:00:00	nsight
0	S	503	14157	14156	2	80	0	-	25208576	futex_	pts/5	00:48:03	java
4	S	0	14621	14605	0	80	0	-	40420	wait	pts/18	00:00:00	su
4	S	0	14632	14621	0	80	0	-	27117	n_tty_	pts/18	00:00:00	bash
0	S	501	16712	16692	0	80	0	-	15047	poll_s	pts/10	00:00:00	ssh
4	S	0	17091	32191	0	80	0	-	40413	wait	pts/22	00:00:00	su
4	S	0	17109	17091	0	80	0	-	27117	n_tty_	pts/22	00:00:00	bash
4	S	0	22140	1	0	80	0	-	27050	wait	pts/22	00:00:00	mysqld_safe
4	S	27	22242	22140	1	80	0	-	144273	poll_s	pts/22	01:16:10	mysqld
0	S	504	23075	17786	0	80	0	-	39382	n_tty_	pts/30	00:00:00	mysql





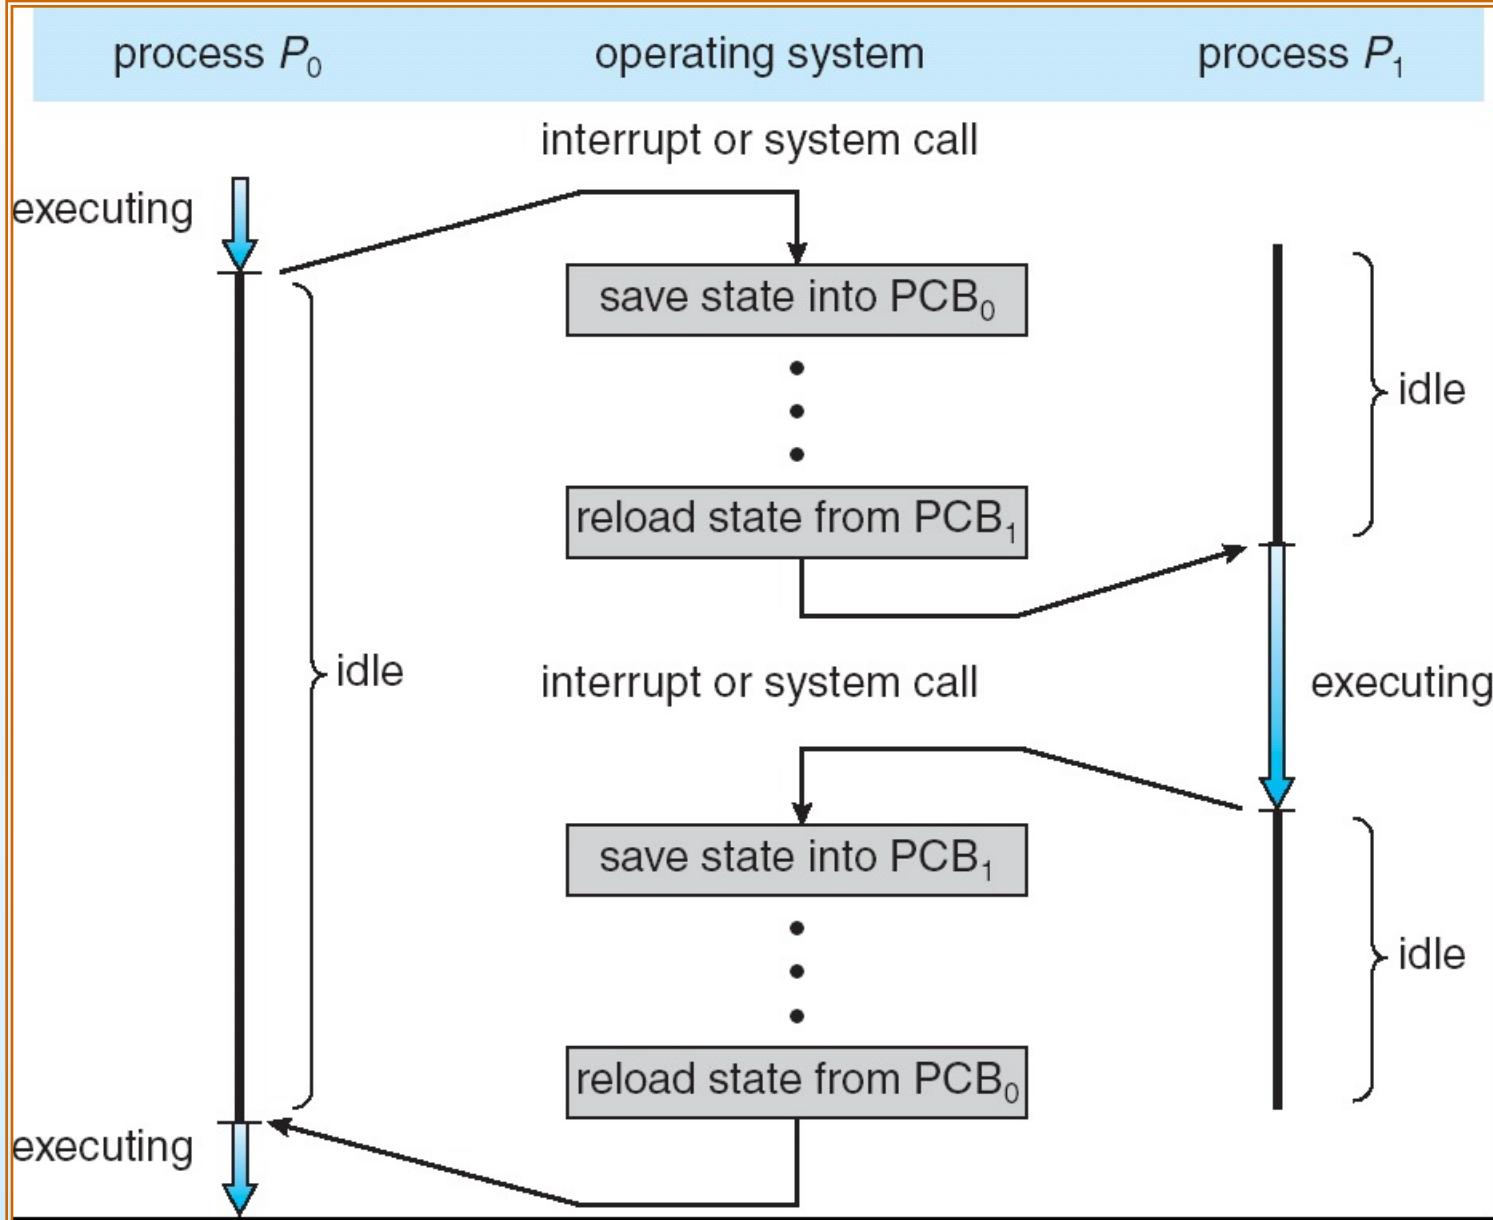
# Process Control Block (PCB)

## Memory area for OS (kernel space)





# CPU Switch From Process to Process





# Chapter 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems

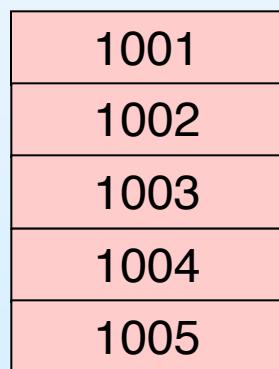
New Teaching Assistant: zhifeipang@zju.edu.cn



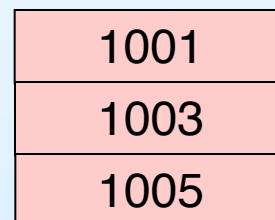


# Process Scheduling Queues

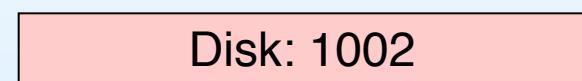
- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- Processes migrate among the various queues



job queue



ready queue



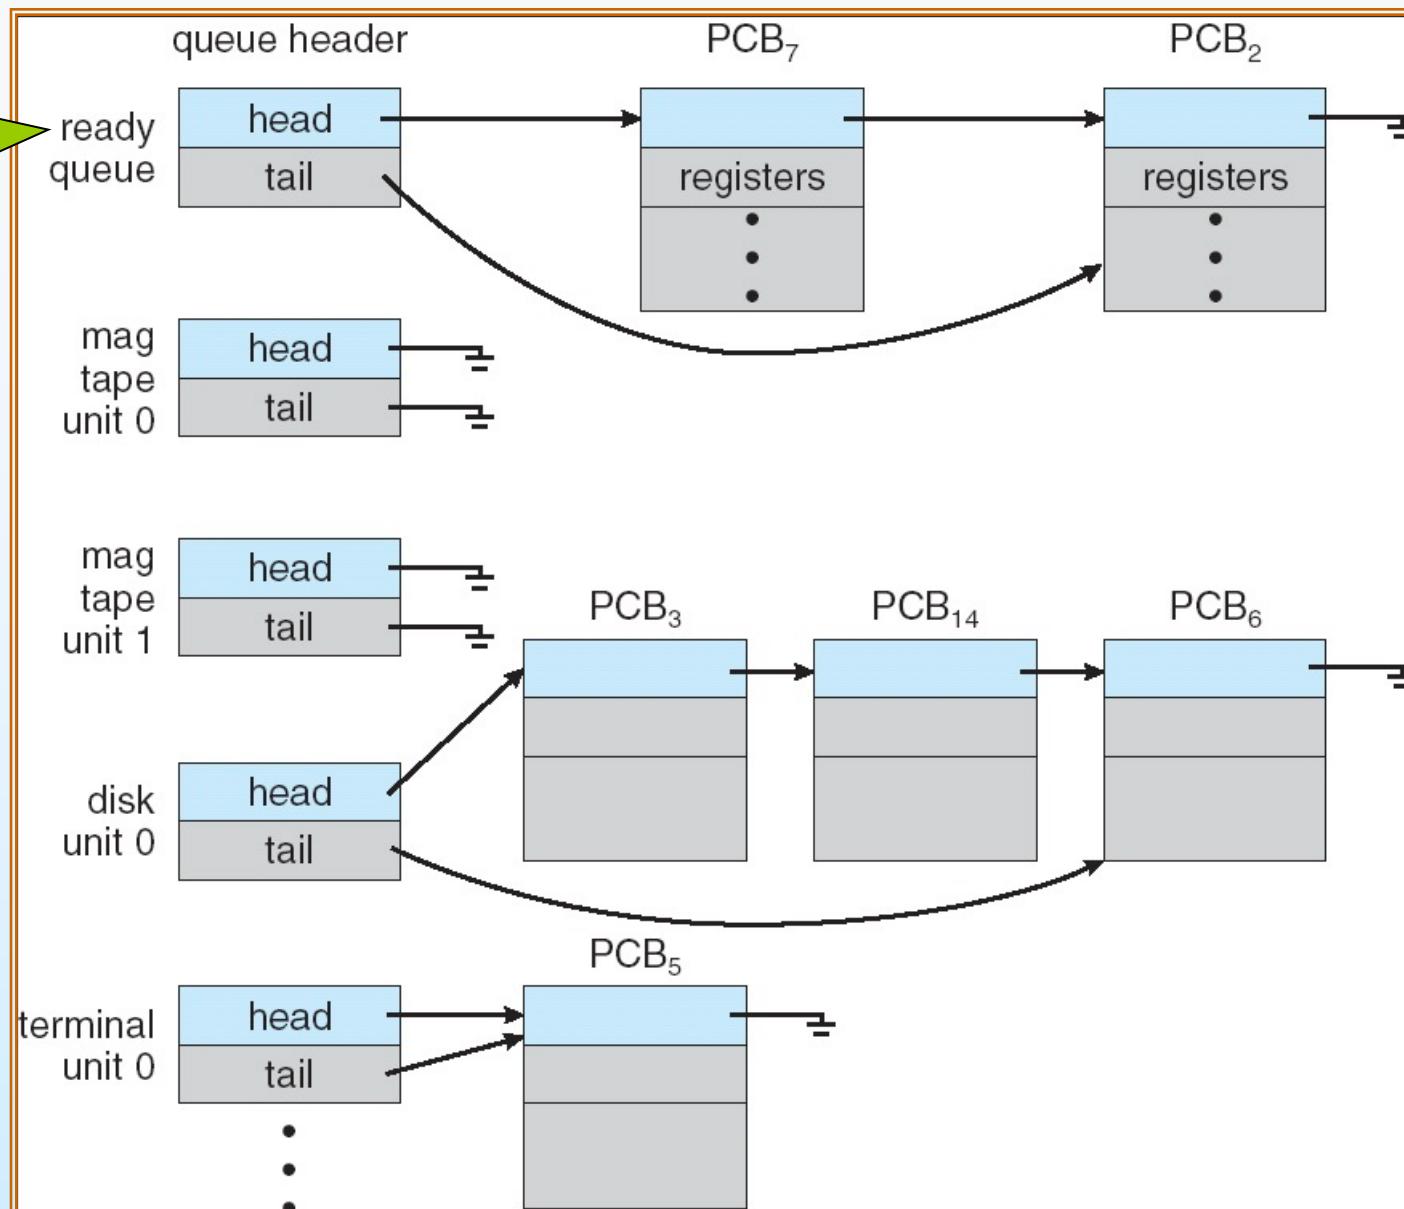
device queue





# Ready Queue And Various I/O Device Queues

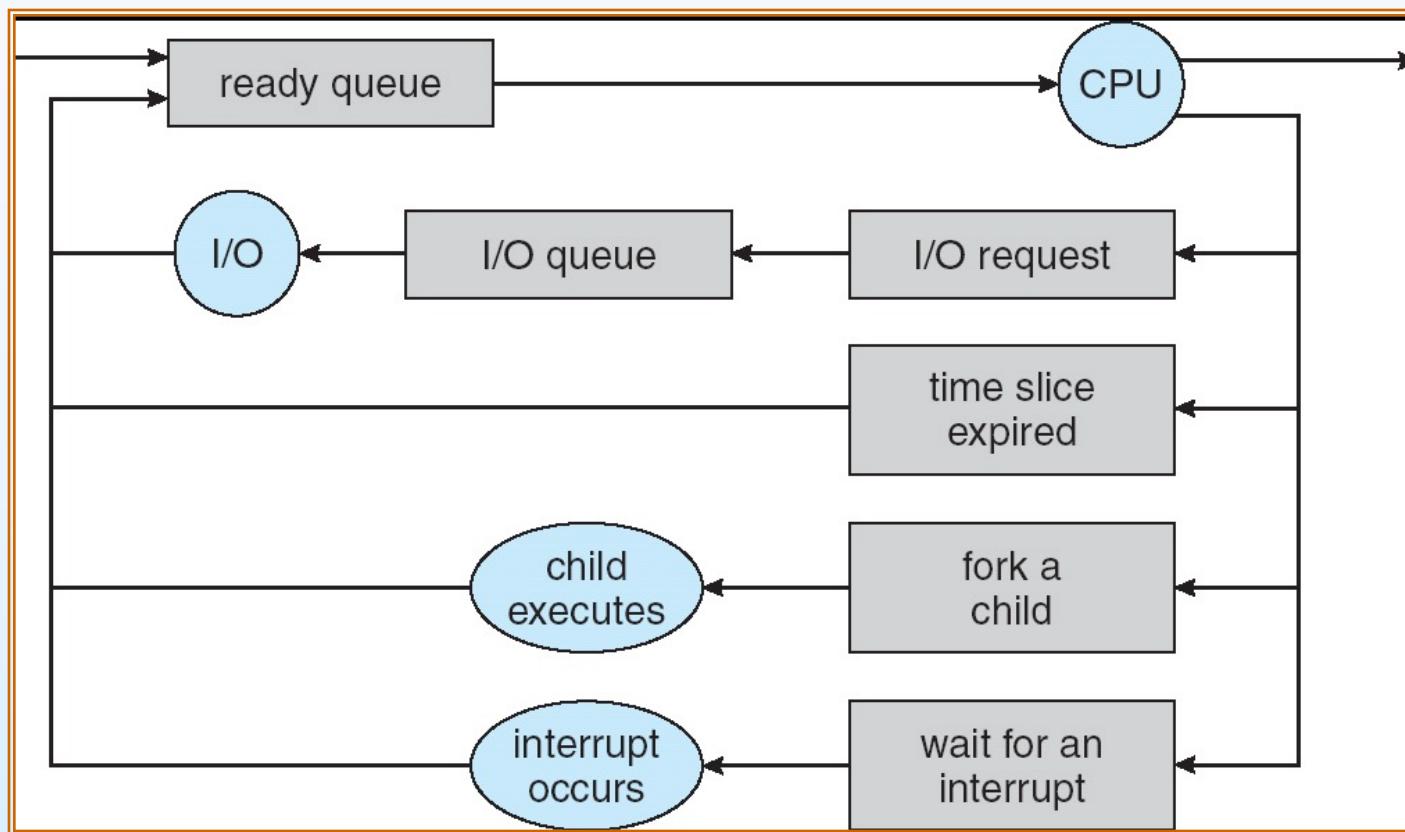
The CPU queue





# Representation of Process Scheduling

A queueing-diagram





# Schedulers

- What is a scheduler? A piece of program
- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into memory (the ready queue)
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU

**UNIX and Windows do not use long-term scheduling**

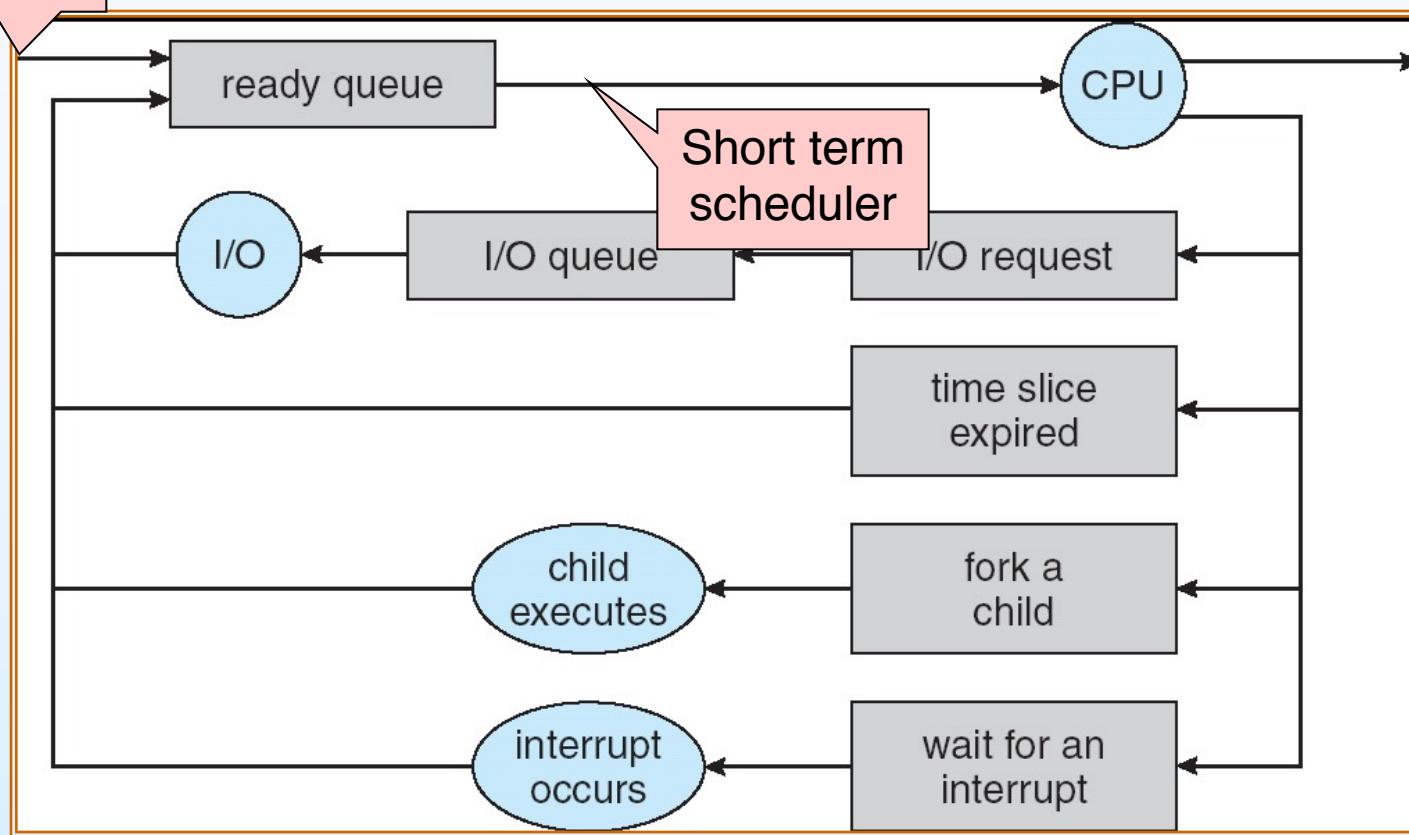




# Representation of Process Scheduling

Long term scheduler

A queueing-diagram





# Representation of Process Scheduling



long term scheduling



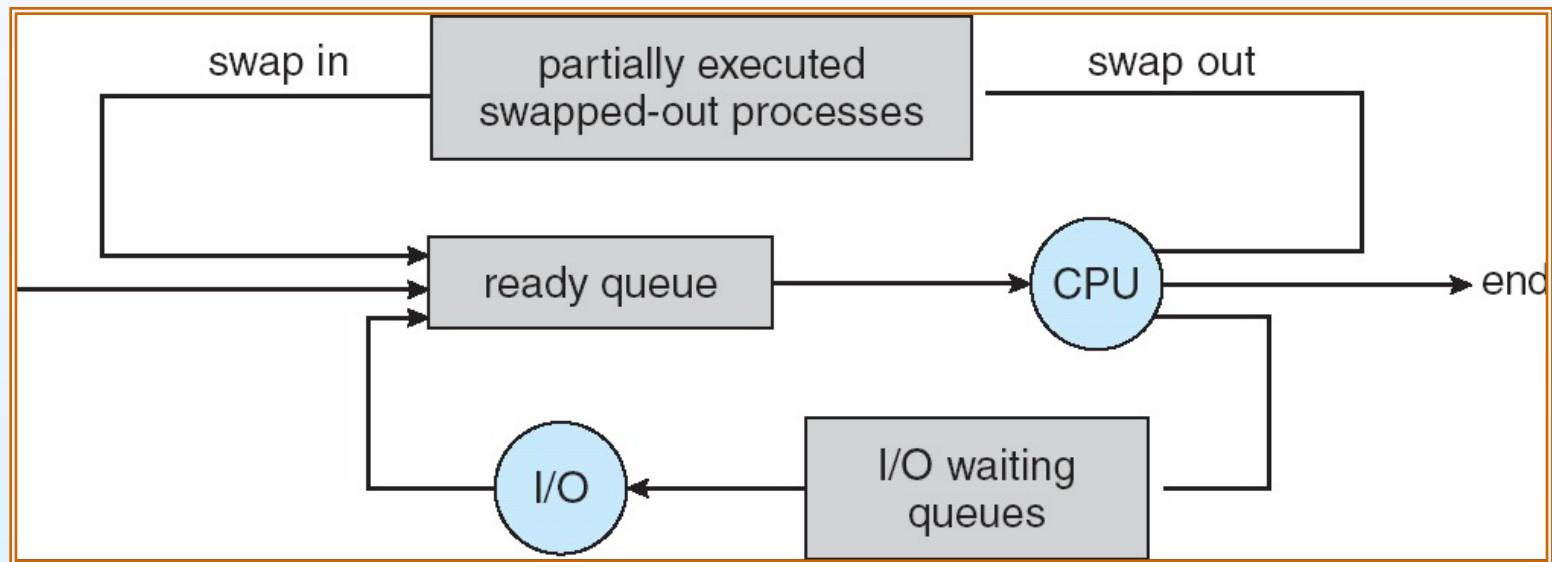
short term scheduling





# Addition of Medium Term Scheduling

Sometimes, it can be good to swap processes out.





# Schedulers (Cont.)

- Short-term scheduler is invoked very frequently (milliseconds) ⇒ (must be fast)
- Long-term scheduler is invoked very infrequently (seconds, minutes) ⇒ (may be slow)
- The long-term scheduler controls the *degree of multiprogramming*
- Processes can be described as either:
  - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
  - **CPU-bound process** – spends more time doing computations; few very long CPU bursts

**Need to select good combination of CPU bound and I/O bound processes.**





# Context Switch

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process
- Context-switch time is overhead; the system does no useful work while switching - typically takes **milliseconds**
- Time dependent on hardware support. In the SPARC architecture, groups of registers are provided.





# Chapter 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems





# Process Creation

- Parent process creates children processes, which, in turn create other processes, forming a tree of processes
- Resource sharing
  - Parent and children share all resources
  - Children share subset of parent's resources
  - Parent and child share no resources
- Execution
  - Parent and children execute concurrently
  - Parent waits until children terminate
- Example of Chrome





# Process Creation (Cont.)

使用入门 × 深入了解 × 浙江大学-首页

www.zju.edu.cn

综合服务 | 办事目录 | 信息公开 | 校网导航 | English

在校生 | 教职工 | 校友 | 考生与访客 | 合作者 | 求职者

校情总览 求是新闻 院系机构 教师队伍 教育教学 科学研究 招生就业 校园生活 搜索

浙江大学 ZHEJIANG UNIVERSITY

新闻头条 视频新闻

星云大师受聘浙大名誉教授

浙大国际海洋土木工程研究中心揭牌

我校代表队荣获全国大学生生物联网设计竞赛...

浙江大学代表队获全国大学生电子设计竞赛...

学会创新、创造、创意

院系部门 公告 学术 文体 交流

扬真我风采，做明日之星--计算机学院与软件...

附属儿童医院与美国洛杉矶儿童医院签署院际合...

化学系彭笑刚教授入选2014年高被引科学家

浙江大学与中科院西安光机所举行共建“浙江大...

浙大出版社4项目获2013年度引进、出版优秀...

浙大出版社获2013中国图书世界影响力出版100强

2014新闻回顾第三编之一：浙大 - 帝国理工日举办学术研讨会

浙江大学-帝国理工学院  
应用数据科学联合实验室  
IMPERIAL COLLEGE AND ZHEJIANG UNIVERSITY  
JOINT LAB FOR APPLIED DATA SCIENCE

> 图书馆 > 校友总会 > 学校邮箱  
> 公开课 > 教育基金会 > 电子地图

浙江省杭州市西湖区余杭塘路866号 | 310058 | 0571-87951111 | 联系方式  
© 2004-2012 浙江大学 | 浙ICP备05074421号 | 宣传部新媒体工作办公室维护

网站地图 | 意见建议 |

微博 微信 网站地图

Operating System Concepts - 7<sup>th</sup> Edition, Feb 7, 2006

3.26

Silberschatz, Galvin and Gagne ©2005





# Process Creation (Cont.)

## ■ Address space

- Child duplicate of parent
- Child has a program loaded into it

## ■ UNIX examples

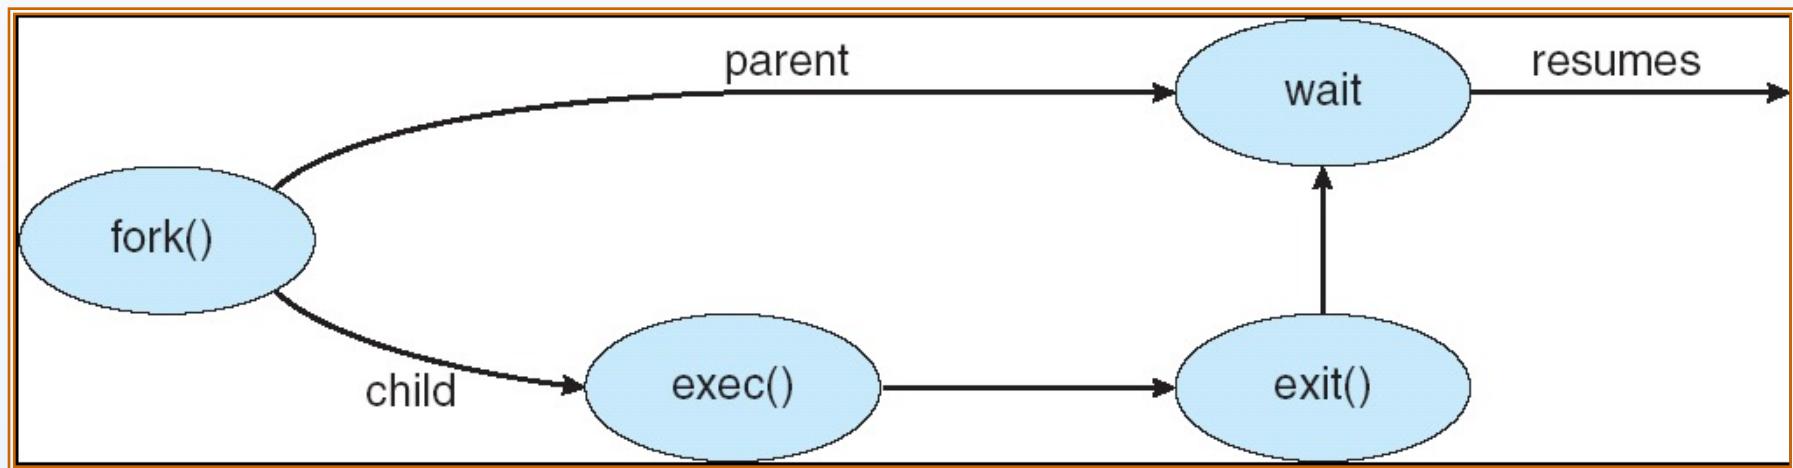
- **fork** system call creates new process
- **exec** system call used after a **fork** to replace the process' memory space with a new program

There are a lot “exec” APIs. For example: *execve()* *execv()* *execle()* *execvp()* *execlp()* etc.





# Process Creation





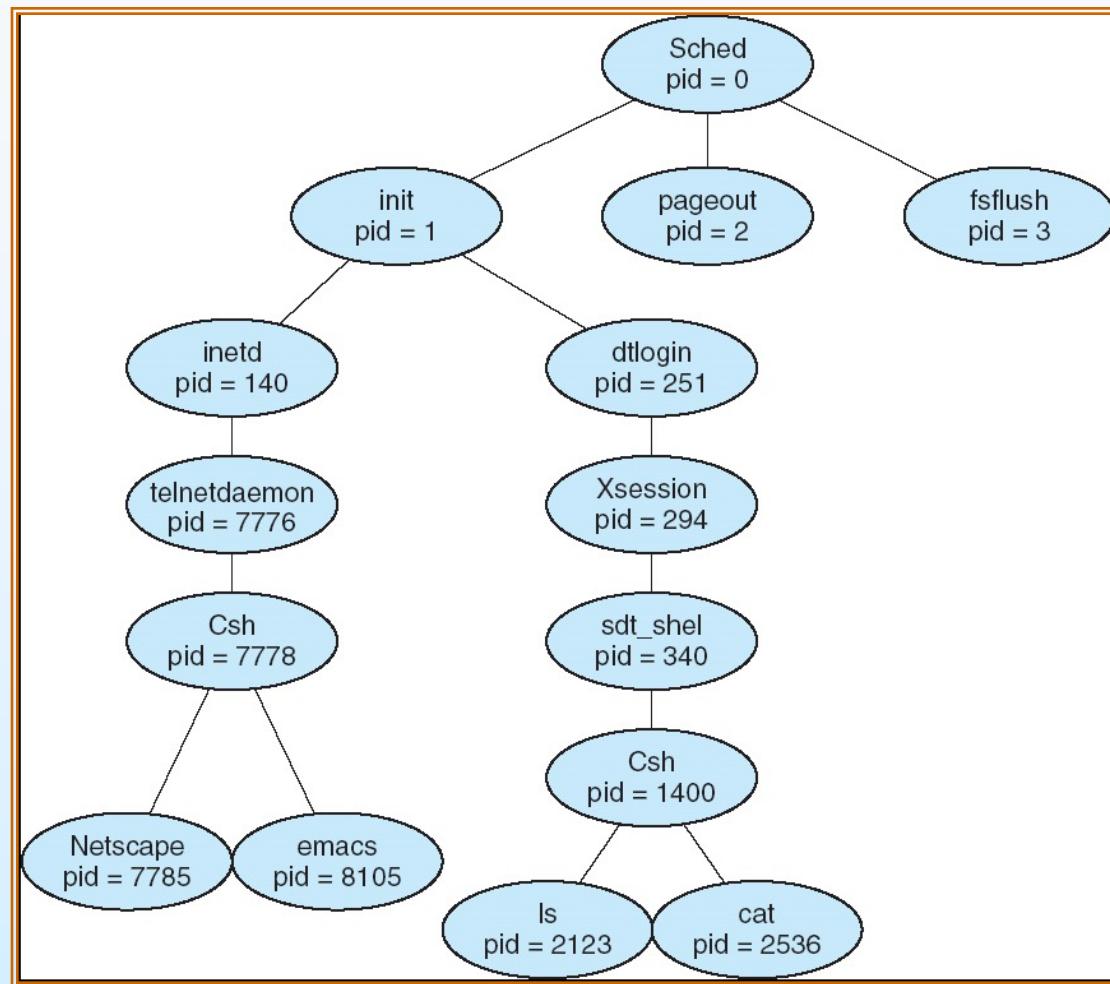
# C Program Forking Separate Process

```
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child
         * to complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```





# A tree of processes on a typical Solaris

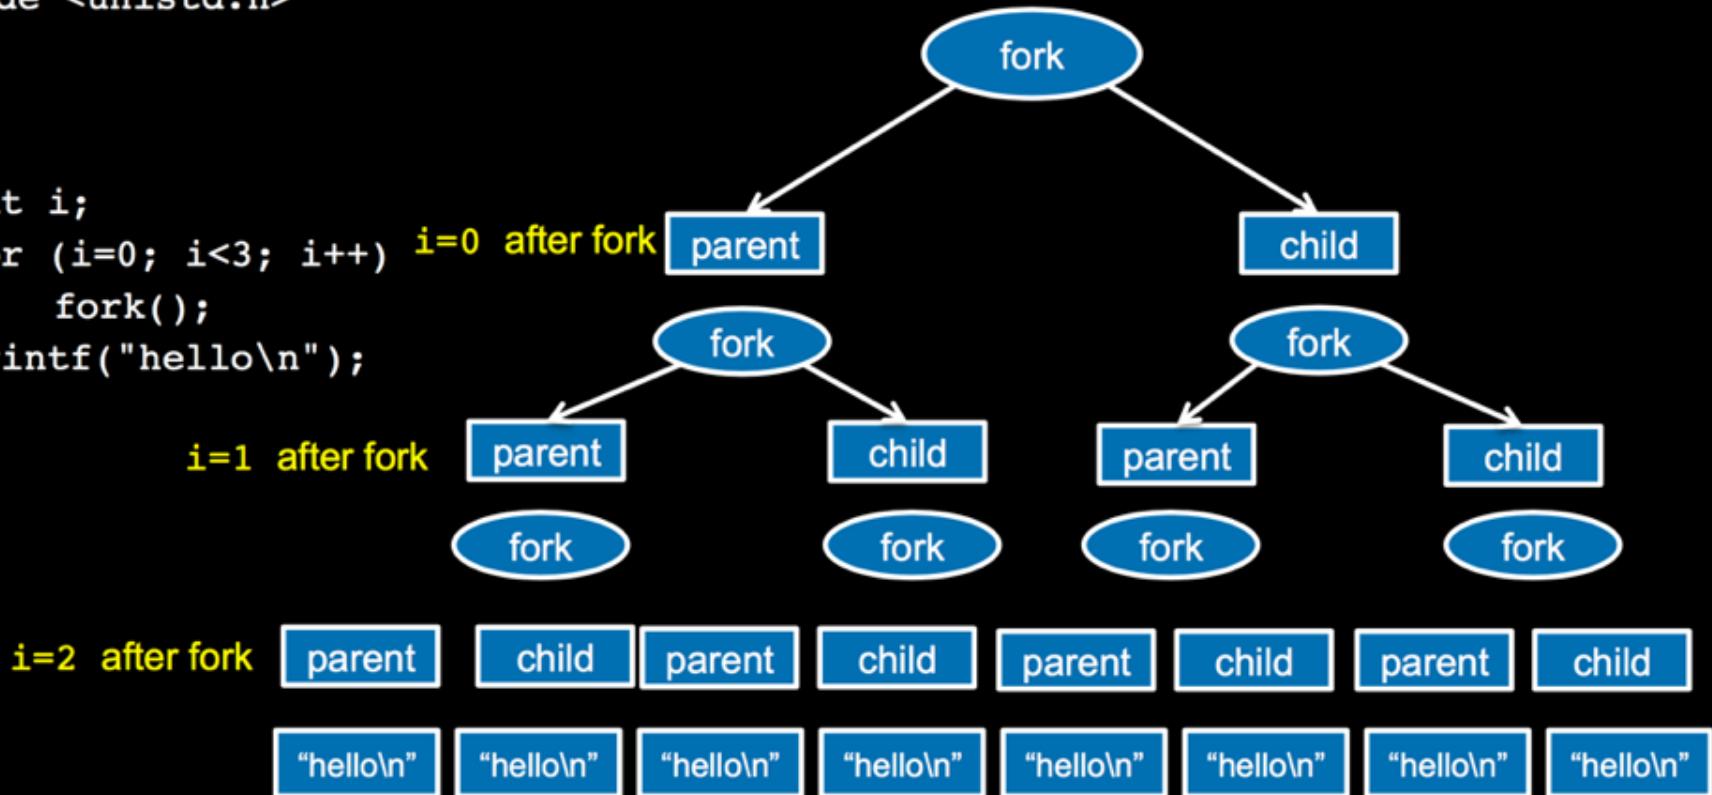




# C Program Forking Separate Process

```
#include <unistd.h>

main()
{
    int i;
    for (i=0; i<3; i++) i=0 after fork
        fork();
    printf("hello\n");
}
```



Answer: 8





# Process Termination

- Process executes last statement and asks the operating system to delete it (**exit**)
  - Output data from child to parent (via **wait**)
  - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (**abort**)
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - If parent is exiting
    - ✓ Some operating systems do not allow child to continue if its parent terminates
      - All children terminated - *cascading termination*
    - ✓ In some other operating systems, the child gets *orphaned*





# Chapter 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- **Cooperating Processes**
- Interprocess Communication
- Communication in Client-Server Systems





# Cooperating Processes

- **Independent** process cannot affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
  - Information sharing
  - Computation speed-up (Multiple CPUs)
  - Modularity
  - Convenience





# Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
  - *unbounded-buffer* places no practical limit on the size of the buffer. Consumer has to wait if no new item.
  - *bounded-buffer* assumes that there is a fixed buffer size. Producer must wait if buffer full.





# Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
```

```
typedef struct {
```

```
    . . .
```

```
} item;
```

```
item buffer[BUFFER_SIZE];
```

```
int in = 0;
```

```
int out = 0;
```

- Solution is correct, but .....





# Bounded-Buffer – Insert() Method

```
while (true) {  
    /* Produce an item */  
    while (((in + 1) % BUFFER_SIZE ==  
out)  
        ; /* do nothing -- no free buffers */  
    buffer[in] = item;  
    in = (in + 1) % BUFFER_SIZE;  
}
```



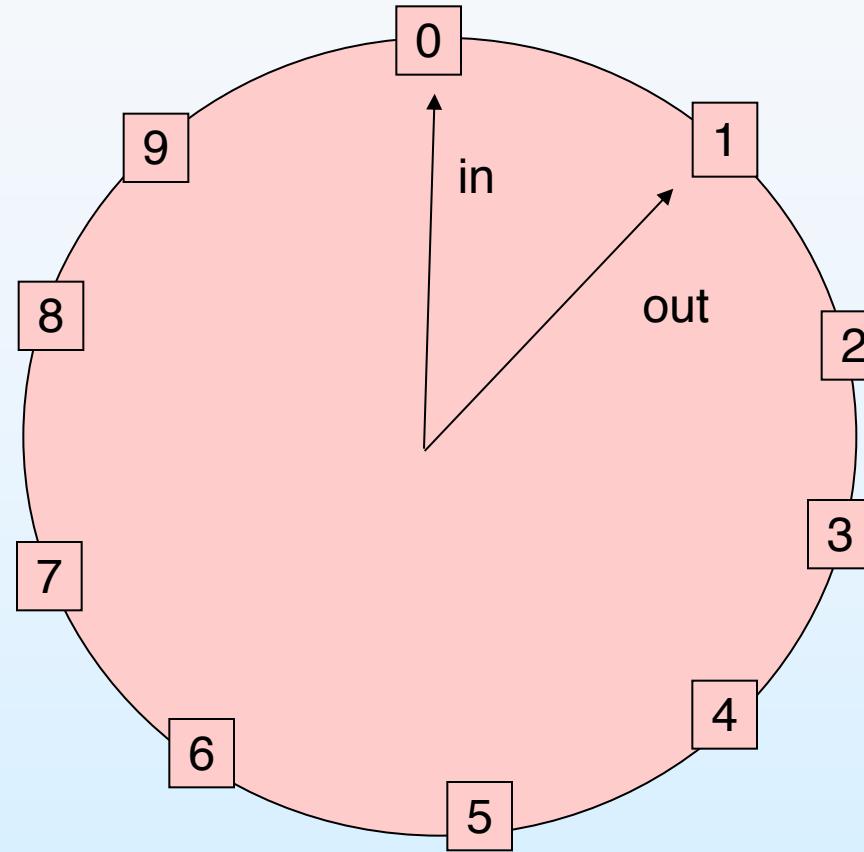


# Bounded Buffer – Remove() Method

```
while (true) {  
    while (in == out)  
        ; // do nothing --  
    nothing to consume  
  
    // remove an item from the  
    buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER SIZE;  
    return item;  
}
```



# Round Table





# Chapter 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- **Interprocess Communication**
- Communication in Client-Server Systems





# Interprocess Communication (IPC)

- Mechanism for processes to communicate and to synchronize their actions
- Models for IPC: *message passing* and *shared memory*
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - **send**(*message*) – message size fixed or variable
  - **receive**(*message*)
- If  $P$  and  $Q$  wish to communicate, they need to:
  - establish a *communication link* between them
  - exchange messages via *send/receive*
- Implementation of communication link
  - physical (e.g., shared memory, hardware bus)
  - logical (e.g., logical properties)





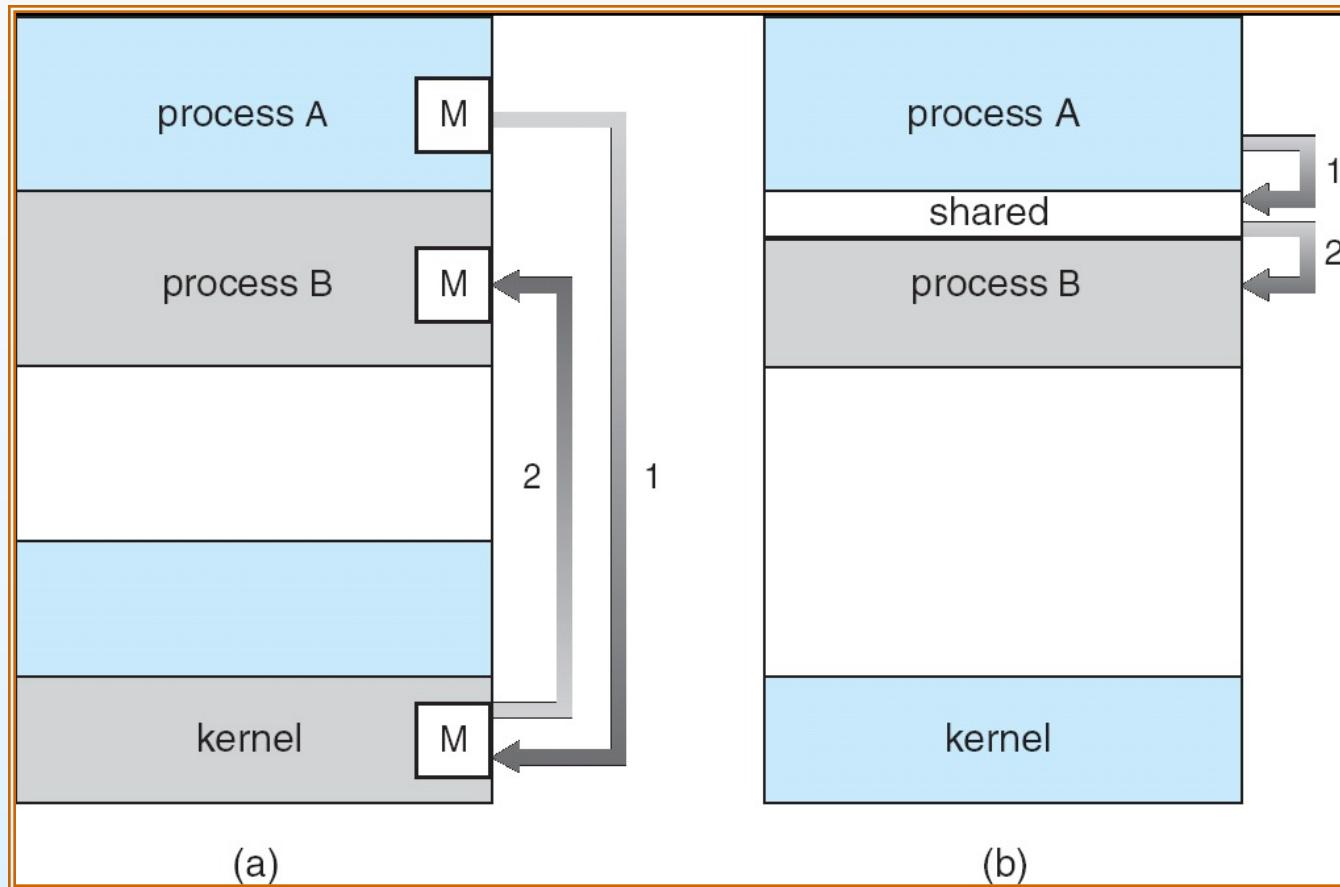
# Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?





# Communications Models



Message passing

Shared memory





# Direct Communication

- Processes must name each other explicitly:
  - **send** ( $P$ , *message*) – send a message to process  $P$
  - **receive**( $Q$ , *message*) – receive a message from process  $Q$
- Properties of communication link
  - Links are established **automatically**
  - A link is associated with **exactly one pair** of communicating processes
  - Between each pair there exists **exactly one link**
  - The link may be unidirectional, but is usually bi-directional





# Indirect Communication

- Messages are directed and received from **mailboxes** (also referred to as **ports**)
  - Each mailbox has a unique id
  - Processes can communicate only if they share a mailbox
- Properties of communication link
  - Link established only if processes share a common mailbox
  - A link may be associated with **many processes**
  - Each pair of processes may share **several communication links**
  - Link may be unidirectional or bi-directional





# Indirect Communication

## ■ Operations

- create a new mailbox
- send and receive messages through mailbox
- destroy a mailbox

## ■ Primitives are defined as:

**send(*A, message*)** – send a message to mailbox A

**receive(*A, message*)** – receive a message from  
mailbox A





# Indirect Communication

## ■ Mailbox sharing

- $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A
- $P_1$ , sends;  $P_2$  and  $P_3$  receive
- Who gets the message?

## ■ Solutions

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver.  
Sender is notified who the receiver was.





# Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
  - **Blocking send** has the sender block until the message is received
  - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
  - **Non-blocking** send has the sender send the message and continue
  - **Non-blocking** receive has the receiver receive a valid message or null





# Buffering

- Queue of messages attached to the link; implemented in one of three ways
  1. Zero capacity – 0 messages  
Sender must wait for receiver (rendezvous)
  2. Bounded capacity – finite length of  $n$  messages  
Sender must wait if link full
  3. Unbounded capacity – infinite length  
Sender never waits
- Control of Buffering





# Chapter 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems





# Client-Server Communication

- Sockets
- Remote Procedure Calls
- Remote Method Invocation (Java)





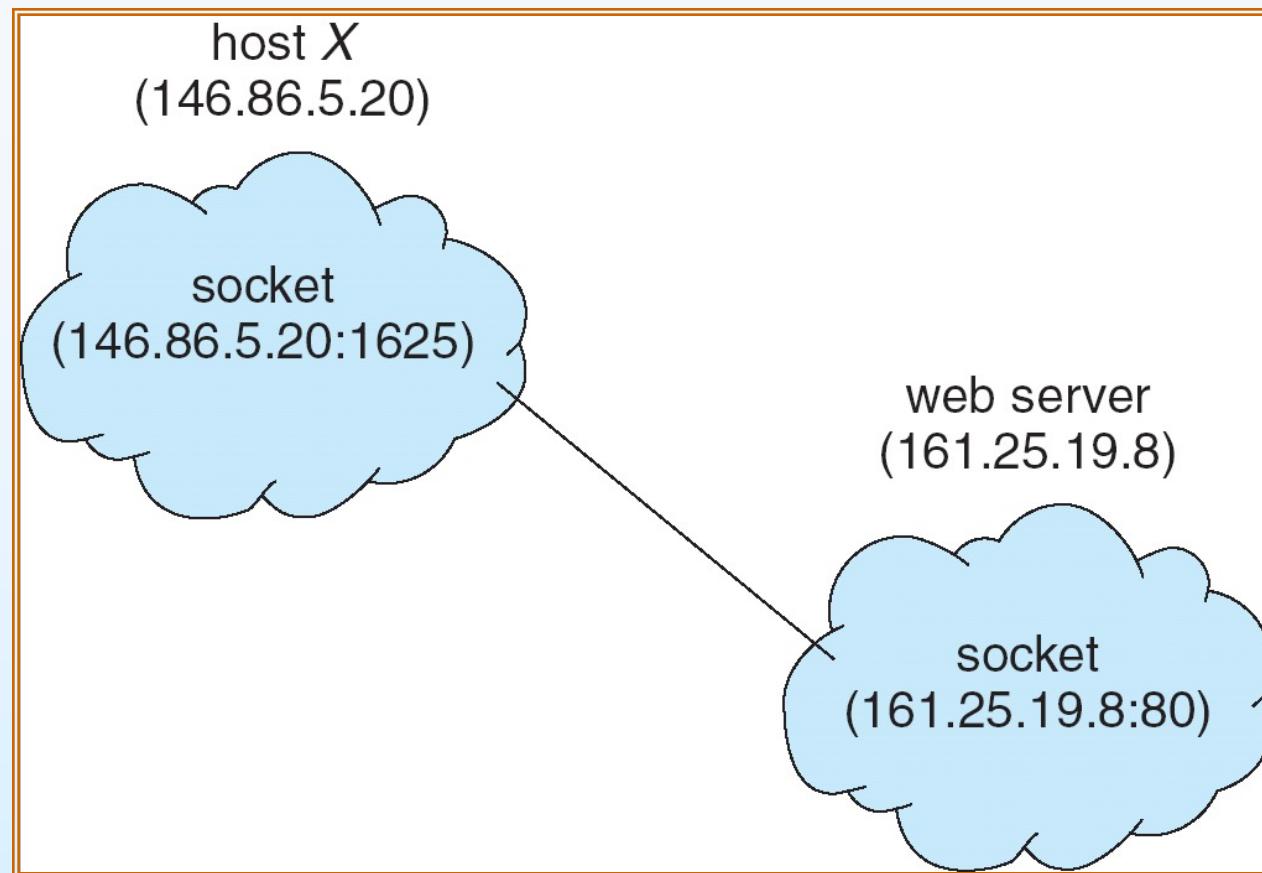
# Sockets

- A socket is defined as an *endpoint for communication*
- Concatenation of IP address and **port**
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets





# Socket Communication





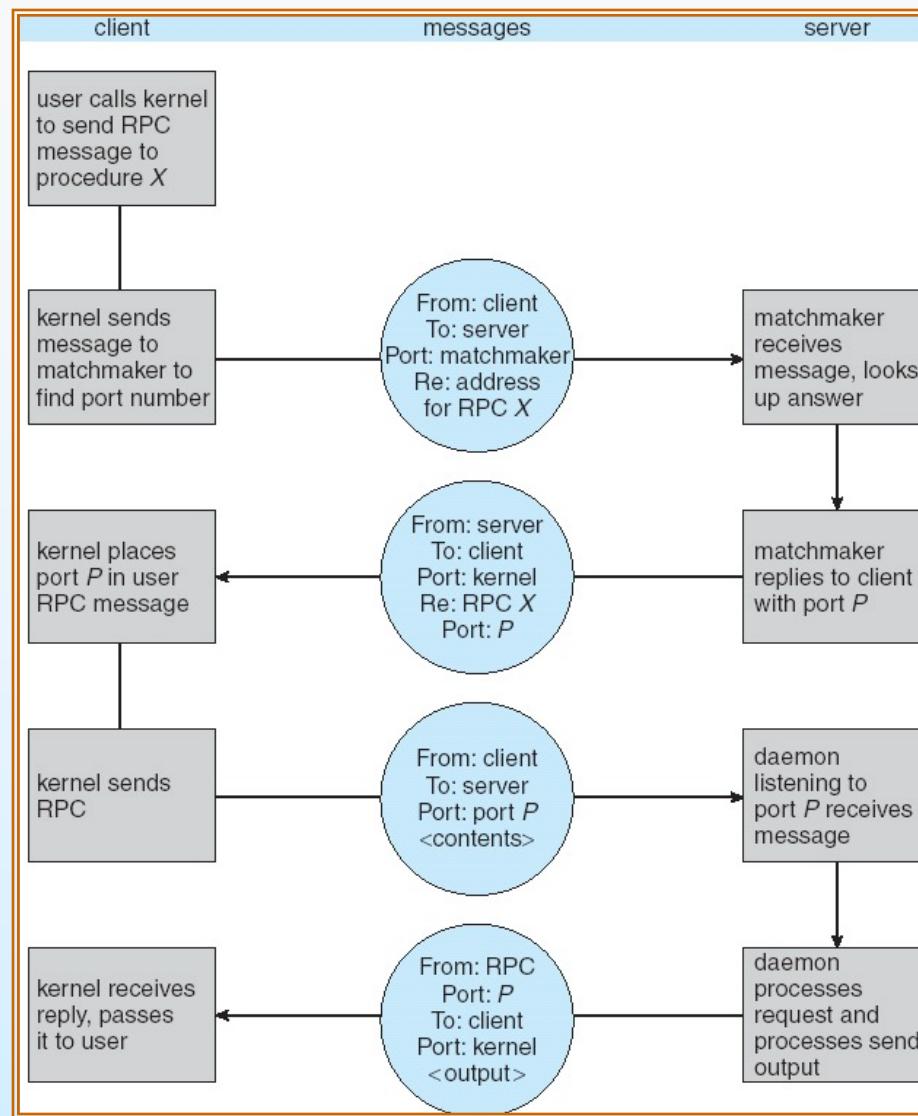
# Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- **Stubs** – client-side proxy for the actual procedure on the server.
- The client-side stub locates the server and *marshals* the parameters.
- The server-side stub receives this message, unpacks the marshaled parameters, and performs the procedure on the server.





# Execution of RPC





# Apache Thrift

## ■ Supported languages:

- [C GLib](#)
- [C++ library](#)
- [C# library](#)
- [D library](#)
- [Erlang library](#)
- [Go library](#)
- [Haskell library](#)
- [Java library](#)
- [JavaScript library](#)
- [node.js library](#)
- [OCaml library](#)
- [Perl library](#)
- [PHP library](#)
- [Python library](#)
- [Ruby library](#)
- [SmallTalk library](#)

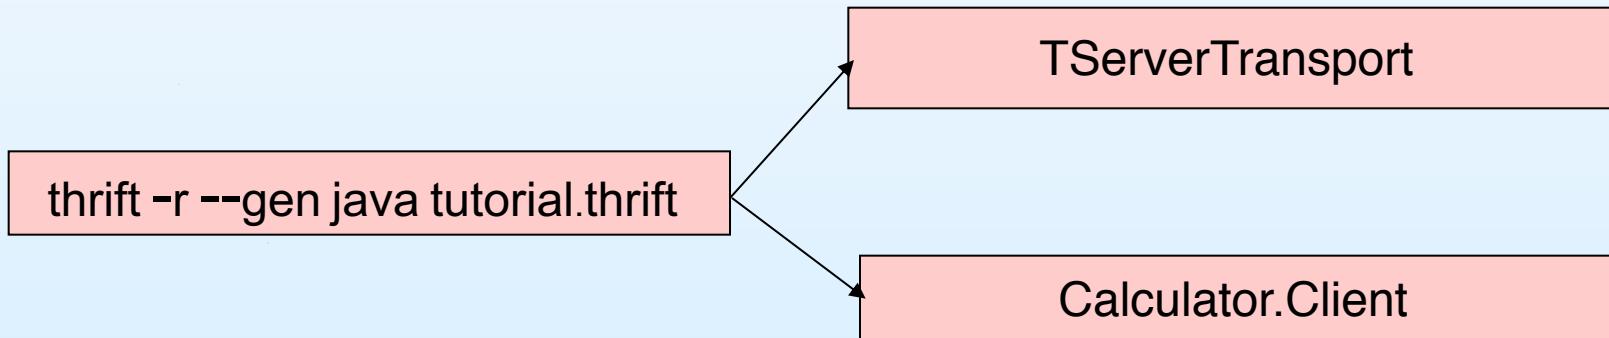
Company	Website
Cloudera	<a href="http://www.cloudera.com">http://www.cloudera.com</a>
Evernote	<a href="http://evernote.com">http://evernote.com</a>
Facebook	<a href="http://www.facebook.com">http://www.facebook.com</a>
last.fm	<a href="http://www.last.fm">http://www.last.fm</a>
Mendeley	<a href="http://www.mendeley.com">http://www.mendeley.com</a>
OpenX	<a href="http://www.openx.org">http://www.openx.org</a>
RapLeaf	<a href="http://www.rapleaf.com">http://www.rapleaf.com</a>
reCaptcha	<a href="http://www.recaptcha.com">http://www.recaptcha.com</a>
Aereo	<a href="http://www.aereo.com">http://www.aereo.com</a>
Siemens	<a href="http://www.siemens.com">http://www.siemens.com</a>





# Apache Thrift

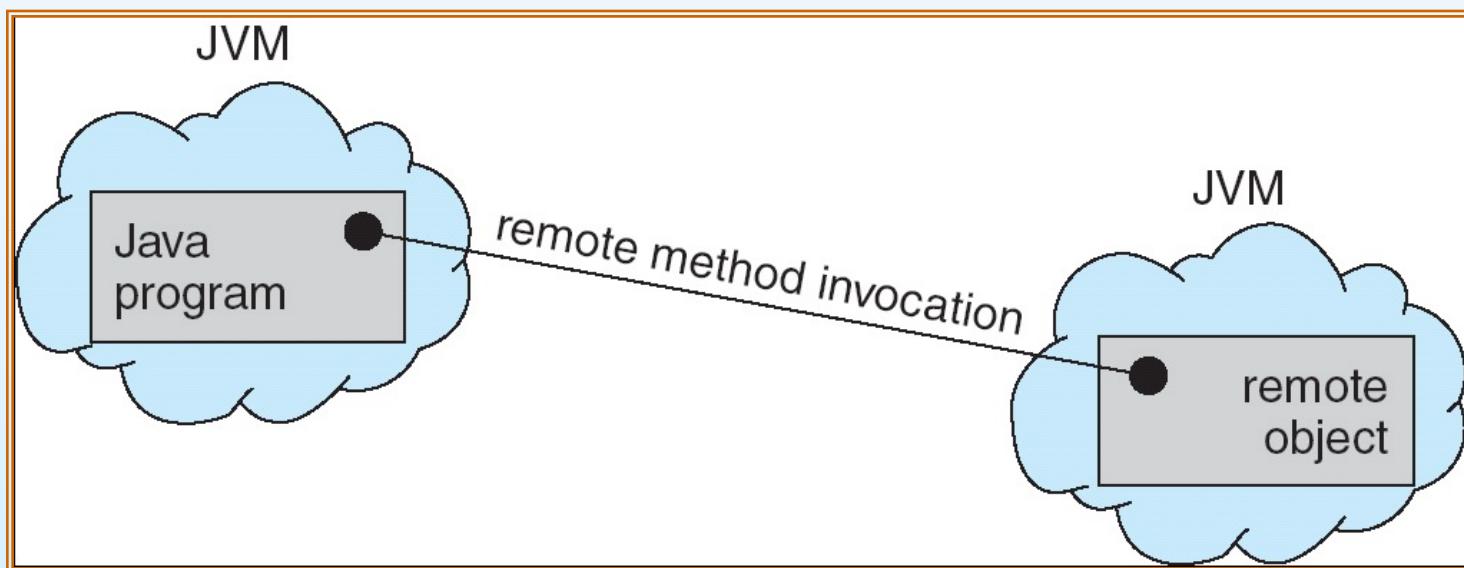
- service Calculator extends shared.SharedService {  
    void ping(),  
    i32 add(1:i32 num1, 2:i32 num2),  
    i32 calculate(1:i32 logid, 2:Work w) throws (1:InvalidOperationException  
                        ouch),  
    void zip() }





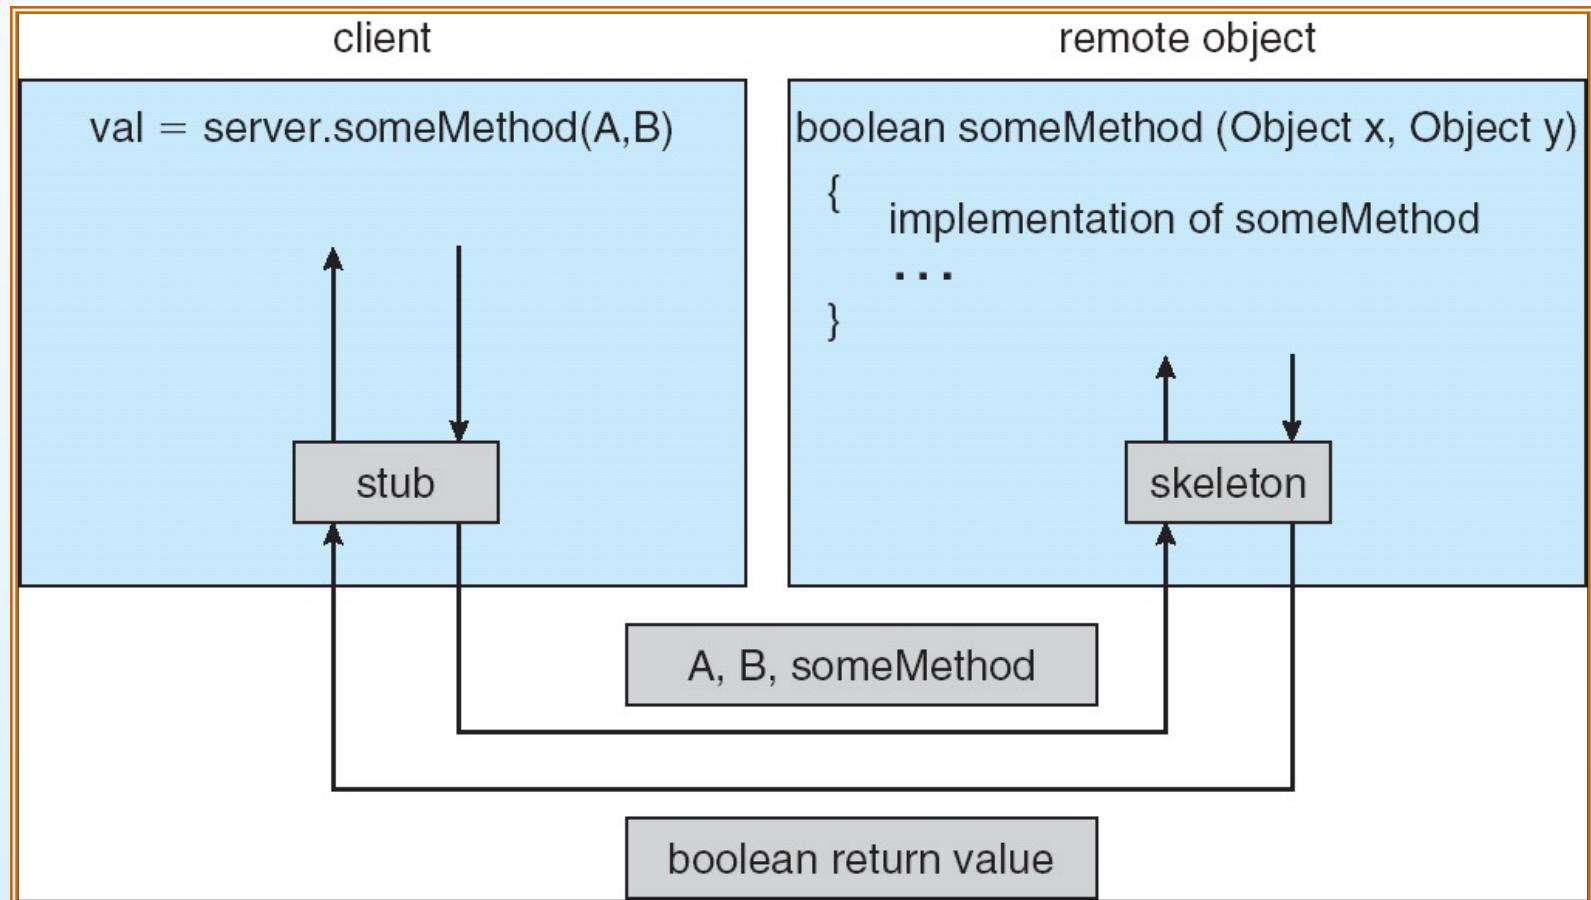
# Remote Method Invocation

- Remote Method Invocation (RMI) is a Java mechanism similar to RPCs.
- RMI allows a Java program on one machine to invoke a method on a remote object.
- RMI is **object-based**





# Marshalling Parameters



Since Java 2 v1.2, skeleton is not needed any more.





The picture can't be displayed.

# End of Chapter 3

